

COMP6771

Advanced C++ Programming

3.2 – Class Types

What is Object-Oriented Programming?

- A class uses data abstraction and encapsulation to define an abstract data type:
 - Abstraction: separation of interface from implementation
 - Useful as class implementation can change over time
 - Encapsulation: enforcement of this via information hiding
- This abstraction leads to two key concepts:
 - Interface: the operations used by the user (an API)
 - Implementation: the data members the bodies of the functions in the interface and any other functions not intended for general use

C++ Class Types

A **class type** in C++ is defined as one of the below three entities:

- **structs**
 - Same as they are in C.
 - A structure made up of data of any type.
- **unions**
 - Same as they are in C.
 - We will not be covering unions much in this course.
- **classes**
 - New in C++.
 - Classes are virtually identical to structs.

C++ Classes

Since you've completed COMP2511 (or equivalent), C++ classes should be straightforward and at a high-level follow very similar principles.

- A class:
 - Defines a new type
 - Is created using the keyword `class`
 - May define some members (methods, data)
 - Contains zero or more public, protected, or private sections
 - Is instantiated through a constructor.
- A member function:
 - must be declared inside the class
 - may be defined inside the class (it is then inline by default)
 - may be declared `const`, when it doesn't modify the data members
- The data members should be private, representing the state of an object.

Classes vs. Structs

- A class and a struct in C++ are almost exactly the same.
- The only technical difference is that:
 - All members of a struct are public by default
 - All members of a class are private by default
- Structs generally are used for simple types with few or no methods.
 - These kinds of structs are called PODs: plain old data.
- Intended use of your type should determine whether it is a struct or class

```
// default everything is private
class foo {
    int member_;
};
```

```
// default everything is public
// we often times denote this with
// a different style of member naming
// note the lack of a trailing _
struct footwo {
    int member;
}
```

Member Access Control

- This is how we support encapsulation and information hiding in C++.
- Can have multiple sections of the same **access specifier**.
- They are:
 - `public` – accessible by everyone
 - `protected` – accessible by children of the class and the class itself.
 - `private` - accessible only by this class
 - `virtual` - cover this in the polymorphism module.

```
class foo {  
    // Members accessible by everyone  
    public:  
        foo(); // The default constructor.  
  
    // Accessible by this class & children.  
    // Will discuss more later in the course  
    protected:  
  
    // Accessible only by this class  
    private:  
        void private_member_function();  
        int private_data_member_;  
  
    // Can define the same access specifier again  
    public:  
};
```

Constructor

- Constructors define how this object should be initialised.
- A constructor has the same name as the class and no return type.
- The constructor that is callable with zero arguments is called the **default constructor**.

```
#include <iostream>

class my_class {
public:
    // this is a constructor
    my_class(int i) {
        i_ = i;
    }
    get_val() {
        return i_;
    }

private:
    int i_;
};

int main() {
    auto mc = my_class{1};
    std::cout << mc.get_val() << "\n";
}
```

Construction Order

Class construction follows this pseudocode algorithm.

```
for each data member in declaration order
  if it has a user-defined initialiser
    Initialise it using the user-defined initialiser
  else if it is a fundamental type (integral, pointer, bool, etc.)
    do nothing (leave it as whatever was in memory before)
  else
    Initialise it using its default constructor
```


Member Initialiser List

- The initialisation phase occurs before the body of the constructor is executed, regardless of whether a member initialiser list is supplied.
- You are able to provide initial values for data members before entering the constructor body.
- This avoids initialising a data member only to have its value overwritten.
- You should **always** use a member initialiser list.
- You also must ensure you initialise data members in declaration order.
 - Otherwise the behaviour is undefined.

```
class user {  
public:  
    user(std::string name, int age)  
        // everything following the colon  
        // is the member initialiser list  
        : name_{name}, age_{age} {  
        /* more complex set-up */  
    }  
  
private:  
    std::string name_;  
    int age_;  
};
```

Delegating Constructor

- A constructor may call another constructor from the member initialiser list.
- Since the other constructor must construct all the data members, you should not specify anything else in the list.
- The other constructor is called completely before this one.
 - This means the lifetime of the object has begun!

```
class point2d {  
public:  
    // the default ctor delegates to the  
    // 2-arg constructor, supplying  
    // default values  
    point2d() : point2d(0, 0) {}  
  
    point2d(float x, float y)  
    : x_{x}, y_{y} {}  
  
private:  
    float x_;  
    float y_;  
};
```

In-class Initialisers

- It can be easy to forget to initialise all data members, especially when there are many.
- You can use in-class initialisers as a last resort fallback if a data member is not initialised with a member initialiser list.
- It is best to minimise their use, as having initialisation logic spread out can be hard to read.

```
class point2d {  
public:  
    point2d(float x, float y)  
        : x_{x}, y_{y} {}  
  
private:  
    // with no default constructor,  
    // the in-class initialisers  
    // will be used instead.  
  
    float x_ = 0;  
    float y_ = 0;  
};
```

explicit

- If a constructor for a class is callable with a single parameter, the compiler will create an implicit type conversion from the parameter to the class type.
- This **may** be the behaviour you want
 - But probably not
- You have to opt-out of this implicit type conversion with the `explicit` keyword.

```
class point2d {
public:
    // single arg parameter
    explicit point2d(float n)
        : x_{n}, y_{n} {}

    // this ctor is also callable with a single arg
    explicit point2d(float *f, int n = 2)
        : x_{f[0]}, y_{n >= 2 ? f[1] : f[0]} {}

private:
    float x_;
    float y_;
};

point2d ex() {
    // error: single arg constructor is explicit
    return 3.0f

    float arr[2] = {0.0f};
    // error: 2-arg constructor also is explicit
    return arr;
}
```

Destructor

- Called when the object goes out of scope.
- How would you use one?
 - Freeing malloc'ed memory.
 - Closing files.
 - Unlocking mutexes.
 - Aborting database transactions.
 - Any kind of resource clean-up.
- `noexcept` is part of C++ exceptions (we will cover this later)

```
// from Java: a boxed integer
class integer {
public:
    integer(int i) : ptr_{new int{i}} {}

    // never forget to free a pointer
    // ever again!
    ~integer() {
        delete ptr_;
    }

private:
    int *ptr_;
}
```

Special Member Functions

- There are six special member functions:
 - The default constructor
 - The copy constructor
 - The move constructor
 - The destructor
 - The copy-assignment operator
 - The move-assignment operator

```
class foo {  
public:  
    // default constructor  
    foo();  
  
    // copy/move (cover later) constructor  
    foo(const foo &other);  
    foo(foo &&other);  
  
    // destructor  
    ~foo();  
  
    // copy/move (cover later) operators  
    foo &operator=(const foo &other);  
    foo &operator(foo &&other);  
};
```

Synthesised Special Members

- The compiler is able to *synthesise* definitions for the special member functions if the user does not provide one.
- The synthesised version:
 - Default construction: tries to default construct all the data members in turn
 - Copy construction: calls the copy constructor of each data member
 - Destruction: goes in *reverse-order*, destructing each data member.
 - Etc.
- It is possible to opt-into synthesis with `default`.
- Likewise, it is possible to opt-out of synthesis with `delete`

```
class point2d {
public:
    // even one user-supplied ctor stops
    // the compiler-generated default.
    point2d(float a, float b);
    point2d() = default; // opt back into it

    // now this class is no longer copyable
    point2d(const point2d &other) = delete;

    // ensure compiler-synthesised dtor.
    ~point2d() = default;

private:
    float x_;
    float y_;
};
```

Removing Unneeded Special Members

- There are several special functions that we must consider when designing classes.
- Ask yourself the question: does it make sense to have this default member function?
 - Yes: Does the compile synthesised function make sense?
 - No: write your own definition
 - Yes: write "<function declaration> = default;"
 - No: write "<function declaration> = delete;"

Non-Static Members

- A class can have member functions (methods).
- A class can also have data members.
- The size of a class only depends on the types and declaration order of its members.
 - Due to alignment requirements, the compiler may insert padding into your class.
- `sizeof(class) >= 1` (why?)

```
class foo {  
public:  
    void speak();  
private:  
    int i;  
    void *ptr;  
    bool b;  
};
```

```
class foo2 {  
public:  
    void speak();  
private:  
    bool b;  
    int i;  
    void *ptr;  
};
```

```
static_assert(sizeof(foo) == 24); // why?  
static_assert(sizeof(foo2) == 16); // why?
```

Incomplete Types

- An incomplete type is a type that has been declared but not defined.
- You can only form pointers and references to incomplete types.
 - Except `void&` is illegal.
- Because of this restriction, a class cannot have data members of its own type.
- Since a class is considered declared once its name has been seen, it can have pointer/reference members to its own type.

```
struct node {  
    int data;  
  
    // node is incomplete.  
    // This is invalid.  
    // This would also make no sense.  
    // What is sizeof(Node)??  
    node next;  
  
    // this, however, is fine.  
    node *next;  
  
    // this is fine, too.  
    node &next;  
};
```

The `this` Pointer

- Member functions have an extra implicit parameter, named `this`.
- This is a pointer to the object on behalf of which the function is called.
- A member function does not explicitly define it, but may explicitly use it
 - As of C++23, you may now define it too.
- The compiler treats an unqualified reference to a class member as being made through the `this` pointer.
- Generally we use a `"_"` suffix for class variables rather than `this->` to identify them
- It is possible to overload a method based on the constness of `this`.

```
class point2d {
public:
    point2d(float x, float y) : x_{x}, y_{y} {}

    float x() {
        // the signature of this method is really
        // float(point2d * const this)
        return x_;
    }

    float x() const {
        // the signature of this method is really
        // float(const point2d * const this)
        return x_;
    }
private:
    float x_ = 0;
    float y_ = 0;
};
```

const Objects

- Member functions are by default only callable by non-const objects.
- You can declare a const member function which is callable by both const and non-const objects
- A const member function may only modify mutable members
 - There are **very few** use cases for mutable
 - One example is as a cache
 - A mutable member *should* mean that the physical state of the member can change without the logical state of the object changing.
 - In practice, this isn't always the case.

```
class point2i {  
public:  
    point2d(int a, int b) : x_{a}, y_{b} {}  
  
    const int& x() const { return x_; }  
    int &x() { return x_; }  
  
    const int& y() const { return y_; }  
    int &y() { return y_; }  
  
private:  
    int x_;  
    int y_;  
}
```

```
const auto p = point2i{1, 2};
```

```
p.x(); // OK! const-qualified method called  
p.y() = 4; // error: calls a non-const method
```

Static Members

- Static members belong to the class, as opposed to any particular object.
- Static methods are callable without any instance.
- Static methods are never const-qualified.
- Static data members' lifetime ends when the program ends.
- Use static members when something is associated with a class, but not a particular instance.

```
class point2d {
public:
    static point2d make_point(float a, float b) {
        return point2d(a, b);
    }

    ~point2d() {
        n_live_ -= 1;
    }

private:
    point2d(float a, float b) : x_{a}, y_{b} {
        n_live_ += 1;
    }

    inline static int n_live_ = 0; // since C++17

    float x_;
    float y_;
};
```

Friends

- A class may declare another function or class as a friend.
 - Note: this does not declare the friend class or function itself.
 - It only says *if* such a class or function exists, then it is a friend.
- Friends are able to access the private members of the class.
- For a class C and friend F:
 - F is not a friend of the children of C.
 - F is not a friend of classes C is a friend of.
 - F is not a friend of any parent classes of C.
- Friends are *always* public.

```
class point2d {
public:
    /* Other implementation details... */
    // declare distance as a friend of point2d
    friend float dist(point2d &, point2d &);

private:
    float x_;
    float y_;
}

float dist(point2d &l, point2d &r) {
    // because dist is a friend, it can access
    // the private members of point2d
    return std::sqrt(l.x_ * r.x_ + l.y_ * r.y_);
};
```

Hidden Friends

- It is possible to declare and define a friend inline in a class.
- This “hidden” friend has different look-up rules than usual:
 - Only discoverable through ADL
 - With a separate declaration outside the class, they are discoverable without ADL.
- Mostly used with operator overloading.

```
namespace hf {
    class point2d {
    public:
        // other implementation details...

        friend float dist(point2d &l, point2d &r) {
            return std::sqrt(l.x_ * r.x_ + l.y_ * r.y_);
        }

    private:
        float x_, y_;
    };

    float dist(point2d &, point2d &);
}

void ex() {
    hf::point2d p = {1, 2}, q = {3, 4};

    // OK: hidden friend found through ADL
    dist(p, q);

    // OK: found through 2nd declaration outside class
    hf::dist(p, q);
}
```

The Power of Friendship

Friendship seems to break encapsulation. Why use it?

- In general, friendship should be used sparingly.
- Friendship can be used between two interconnected classes.
 - E.g., a container and its iterator.
- Friendship can also be used for **API Extension**.
 - This does not mean unrelated code can extend an API.
 - It is for a different calling style for a logical operation of a type.
 - E.g., for a `vec3` class, the `cross()` operation should be a friend since `cross(a, b)` makes more sense than `a.cross(b)`.
- For certain operator overloads, it is more convenient for them to be defined as non-member friend functions.

Interface & Implementation

- We are used to declaring functions in header files and defining them in .cpp files.
- With classes we have two options:
 - Define the class completely in the header.
 - Define the class in the header but only declare methods. Define methods in the .cpp file.
- Defining everything in the header is easier for us but can slow down compilation due to duplicate code.
- Either way, static data members must be defined in a .cpp file unless declared inline.

```
// point.hpp - interface file
class point2d {
public:
    point2d(float a = 0, float b = 0);
```

```
    float &x() { return x_; }
```

```
private:
    static int cnt;
    float x_;
    float y_;
};
```

```
// point.cpp - implementation file
int point2d::cnt = 0; // static member
```

```
point2d::point2d(float a, float b)
: x_{a}, y_{b} {} // ctor implementation
```

Feedback (stop recording)

