

# COMP6771

## Advanced C++ Programming

### 5.2 – Resource Management

# In this Lecture

## Why?

- Performance & Control
  - with great power comes great responsibility.
- Leak freedom in C++.
- C++ is not garbage collected.
- While we have ignored heap resources (malloc/free) to date, they are a critical part of many libraries, and we need to understand best practices around usage.

## What?

- Resource Management
- RAI
- Smart Pointers

# Revision: Objects

- What is an object in C++?
  - An object is a region of memory associated with a type.
  - Unlike some other languages (Java), basic types such as `int` and `bool` are objects.
- For the most part, C++ objects are designed to be intuitive to use.
- Special things can we do with objects:
  - Create them;
  - Destroy them;
  - Copy them; and
  - Move them.

# Long Lifetimes

- There are 3 ways you can try and make an object in C++ have a lifetime that outlives the scope it was defined in:
  - Returning it out of a function via copy (can have limitations).
  - Returning it out of a function via references (leads to dangling references).
  - Returning it out of a function as a heap resource (explicit memory management).

```
struct point2i { int x; int y;}

// This function returns a new object,
// Not a reference to the object
point2i make_point_copy(int x, int y) {
    point2i p = {x, y};
    return p;
}

// Returns a reference to a stack-local variable.
// Program has undefined behaviour if used.
point2i &make_point_dangling(int x, int y) {
    point2i p = {x, y};
    return p;
}

// Returns a reference to a heap-allocated variable.
// Now we must manage the memory explicitly.
point2i &make_point_heap(int x, int y) {
    point2i *ptr = new point2i{x, y};
    return *ptr;
}
```

# Long Lifetimes & References

- We need to be very careful when returning references.
- *The object **must** always outlive the reference.*
- Using dangling references is undefined behaviour!
- **Moral of the story:** Do not return references to stack-local variables.
- Unless we copy, for objects we create **INSIDE** a function, we must heap-allocate them!

```
auto okay(int &i) -> int& {  
    return i;  
}
```

```
// why is this ok?  
// answer: lifetime extension  
auto okay(int &i) -> const int& {  
    return i;  
}
```

```
auto not_okay(int i) -> int& {  
    return i;  
}
```

```
auto definitely_not_okay() -> int& {  
    auto i = 0;  
    return i;  
}
```

# new & delete

- Non-global objects are either stored on the *stack or heap*.
- In general, you've created most objects on the stack (without realising!)
- We can allocate heap objects via `new` and free them via `delete`.
  - `new` and `delete` call the constructors/destructors of what they are creating.

```
#include <cstdlib>
#include <iostream>
#include <vector>

int main() {
    int *a = new int{4};
    auto *b = new std::vector<int>{1,2,3};

    std::cout << *a << "\n";
    std::cout << (*b)[0] << "\n";

    delete a; // frees a
    delete b; // frees b

    return EXIT_SUCCESS;
}
```

# A Heap of Memory

- Why do we need heap resources?
  - Heap objects can outlive the scopes they were created in.
  - More useful in contexts where we need explicit control of ongoing memory size e.g. for vector.
  - Stack has limited space on it for storage, heap is much larger.
  - No matter how much we try, it is very difficult to use all of dynamically allocated memory.
    - Primary exception is on embedded systems.

```
#include <iostream>
#include <vector>

int *make_int(int i) {
    int *a = new int{i};
    return a;
}

int main() {
    int *a = make_int(6771);

    // a was defined in a scope that
    // no longer exists
    std::cout << *a << "\n";

    delete a;
}
```

# Heap Allocation is Not Free

- Non-trivial programs almost always use the heap in some way.
  - All of our examples using `std::vector` have been using the heap secretly as well.
- Heap allocation is expensive compared to stack allocation.
  - Should be avoided in performance-critical code.
- C++ has value semantics.
  - Is it possible to avoid the cost associated with deep-copying a heap allocated object?
- Answer: Move Semantics!



# Move Semantics (C++11 onwards)

- Move semantics are a way for class-types to “steal” the data of a temporary value and reuse it.
- Makes use of a language feature called **rvalue references**.
- Let’s review lvalues and rvalues.
  - Also called “value categories”.

```
#include <vector>

std::vector<int> make_vec() {
    return std::vector<int>{1, 2, 3, 4, 5};
}

int main() {
    // here, v is default initialised.
    std::vector<int> v;

    // The vector from make_vec() is a temporary.
    // Rather than make v allocate a vector that
    // can fit the contents of make_vec(), why not
    // just steal make_vec()'s data directly?
    // It is going out of scope, so nothing will be
    // affected by this.
    v = make_vec(); // make_vec() is “moved” into v.
}
```

# lvalue vs rvalue

- **lvalue**: An expression that is an object reference.
  - E.g. named variables.
  - You can always take the address of an lvalue.
- **rvalue**: Expression that is not an lvalue
  - E.g.. Object literals, return result of functions.
  - rvalues are temporary and short lived, while lvalues live a longer life since they exist as variables.

```
int &f();

int main() {
    // 5 is rvalue, i is lvalue
    int i = 5;

    // j is lvalue, i is lvalue
    int j = i;

    // 4 + i produces rvalue then stored in lvalue k
    int k = 4 + i;

    // error: cannot assign to rvalues.
    6 = k;

    // taking address of an lvalue
    int* y = &k;

    // error: cannot take address of an rvalue.
    int* y = &6771;

    // OK: f() returns an lvalue reference
    f() = 3;
}
```

# Ivalue References

- There are multiple types of references:
  - lvalue references look like T&.
  - lvalue references to const look like const T&.
- Once the lvalue reference goes out of scope, the original lvalue is still usable.
- Constant lvalue references exhibit **lifetime extension**.
  - Any temporaries bound to a const T& act like an lvalue.
  - The temporary gets stored in memory and is referred to via the reference

```
// regular lvalue  
int y = 10;
```

```
// OK: binding lvalue to an lvalue reference.  
int &yref = y;
```

```
// OK: binding lvalue to a const lvalue  
reference.  
const int &cref = y;
```

```
// !!: cannot bind const lvalue to a non-const  
ref  
int &ref = static_cast<const int>(y);
```

```
// !!: rvalues only bind to const lvalue  
references  
const int &extended = 10;
```

# rvalue References

- rvalue references look like T&&.
- rvalue references extend the lifespan of the temporary object to which they are assigned.
- Non-const rvalue references allow you to modify the rvalue.
- An rvalue reference formal parameter means that the value was disposable from the caller of the function.
  - If the callee modified value, who would notice or care?
  - The caller has promised that it won't be used anymore
  - An rvalue reference parameter is an lvalue inside the function.

```
#include <iostream>
```

```
// Declaring an rvalue reference  
int &&rref = 20;
```

```
void inner(std::string &&value) {  
    value[0] = 'H';  
    std::cout << value << '\n';  
}
```

```
void outer(std::string &&value) {  
    inner(value); // This call fails? Why?  
    std::cout << value << '\n';  
}
```

```
int main() {  
    outer("hello"); // This call succeeds.  
    auto s = std::string("hello");  
  
    // This call fails because s is an lvalue.  
    inner(s);  
}
```

# std::move

## Uses of rvalue references:

- They are used in working with the move constructor and move assignment special member functions.
- cannot bind non-const lvalue reference of type 'int&' to an rvalue of type 'int'.
- cannot bind rvalue references of type 'int&&' to lvalue of type 'int'.
- A library function that converts an lvalue to an rvalue so that a "move constructor" (similar to copy constructor) can use it.
- This says "I don't care about this anymore".
- All this does is allow the compiler to use rvalue reference overloads.

```
// std::move looks something like this.  
T&& move(T& value) {  
    return static_cast<T&&>(value);  
}
```

```
void inner(std::string&& value) {  
    value[0] = 'H';  
    std::cout << value << '\n';  
}
```

```
void outer(std::string &&value) {  
    // this now works!  
    inner(std::move(value));  
  
    // Value is now in a valid but unspecified state.  
    // Although this isn't an error, this is bad code.  
    // Only access moved variables to remake them.  
    std::cout << value << '\n';  
}
```

```
int main() {  
    outer("hello"); // This works fine.  
    auto s = std::string("hello");  
    inner(s); // This fails because i is an lvalue.  
}
```

# RAII: Resource Acquisition is Initialisation


- A resource is anything that is scarce and must be managed.
  - E.g., Memory, locks, file descriptors, time.
- RAII is a concept where we encapsulate resources inside objects.
  - Acquire the resource in the constructor.
  - Release the resource in the destructor.
  - Deep-copy the resource if it makes sense.
  - Transfer ownership if it makes sense.
  - Resource is always released at a known point in the program, which you can control.
- Every resource should be owned by either:
  - Another resource (e.g. smart pointer, data member).
  - Named resource on the stack.
  - A nameless temporary variable.

# RAII-aware Classes

- When writing a class, if we can't default all of our operators (preferred), we should consider the **Rule of 5**:
  1. Destructor
  2. Move assignment
  3. Move constructor
  4. Copy constructor
  5. Copy assignment
- The presence or absence of these 5 operations is critical in managing resources.
- If you define one of these special member functions, you should explicitly define or default all of them.

# Implementing an RAI Class

- Let's try implementing our own version of `std::vector`.
- It's going to have to manage some heap memory, so it should look something like this.



```
class vec {
public:
    vec(int n)
        : ptr_{new double[n]}, size_{n}, cap_{n} {}

    vec(const vec &other);
    vec(vec &&other);

    vec &operator=(const vec &other);
    vec &operator=(vec &&other) noexcept;

    ~vec();

    // other implementation...
private:
    double *ptr_;
    int size_;
    int cap_;
};
```



# RAII Class: Destructor

- What happens when `v` goes out of scope?
  - Destructors are called on each member...
  - But destructing a pointer type does nothing!
- As it stands, this will result in a memory leak.
  - How do we fix?
  - Use the destructor!

```
class vec {  
public:  
    // other implementation...  
    ~vec() noexcept /* optional noexcept specifier */ {  
        delete[] ptr_;  
    }  
private:  
    double *ptr_;  
    int size_;  
    int cap_;  
};  
  
int main() {  
    auto v = vec{4};  
} // went out of scope...problem?
```

# RAII Class: Move Constructor

- Move constructor should be noexcept.
- Can use `std::exchange` to exchange the important data members from the temporary.
- Often not much more complex than this.

```
class vec {  
public:  
    // other implementation...  
    vec(vec &&other) noexcept  
        : data_{std::exchange(other.data_, nullptr)},  
          size_{std::exchange(other.size_, 0)},  
          cap_{std::exchange(other.cap_, 0)} {}  
  
private:  
    double *ptr_;  
    int size_;  
    int cap_;  
};
```

# RAII Class: Move Assignment

- Like the Move Constructor, but the destination has already been constructed.
- The easiest way to write a move assignment is generally to do member-wise swaps, then clean up the original object.
- Doing so may mean some redundant code, but it means you don't need to deal with mixed state between objects.

```
class vec {  
vec &operator=(vec &&other) noexcept {  
    if (this != &other) {  
        std::swap(data_, other.data_);  
        std::swap(size_, other.size_);  
        std::swap(cap_, other.cap_);  
  
        delete[] other.data_;  
        other.data_ = nullptr;  
        other.size_ = 0;  
        other.capacity = 0;  
    }  
}  
// other implementation...  
};
```

# Moving Objects

- Always declare your moves as `noexcept` (why?)
  - Failing to do so *might* make your code slower.
- Unless otherwise specified, objects that have been moved from are in a valid but unspecified state.
- Moving is an optimisation on copying
  - The only difference is that when moving, the moved-from object is mutable.
  - Not all types can take advantage of this:
    - If moving an `int`, mutating the moved-from `int` is extra work
    - If moving a `vector`, mutating the moved-from `vector` potentially saves a lot of work.
- Moved from objects must be placed in a valid state.
  - Moved-from containers usually contain the default-constructed value.
  - Moved-from types that are cheap to copy are usually unmodified.
  - Although this is the only requirement, individual types may add their own constraints.
- Compiler-generated move constructor / assignment performs member-wise moves.

# RAII Class: Copy Constructor

- What does the default synthesised copy constructor do?
  - Member-wise copy.
- What are the consequences?
  - Any modification to our vec instance will also change the original.
  - Likely lead to a double-free.
- How can we fix this?
  - Custom copy constructor that does a deep copy!
- Copy constructors are **rarely** noexcept (why?).

```
class vec {  
public:  
    // other implementation...  
    vec(const vec &o)  
        : data_{new double[o.size_]},  
          size_{o.size_},  
          cap_{o.cap_} {  
        std::copy(o.data_, o.data_ + o.size_, data_);  
    }  
  
private:  
    double *ptr_;  
    int size_;  
    int cap_;  
};
```

# RAII Class: Copy Assignment

- Assignment is the same as construction, except there is already a constructed object in your destination.
- You need to clean up the destination first.
- The copy-and-swap idiom makes this trivial.

```
class vec {  
public:  
    // other implementation...  
    vec& operator=(const vec &other) {  
        vec(other).swap(*this);  
        return *this;  
    }  
private:  
    void swap(vec &other) {  
        std::swap(data_, other.data_);  
        std::swap(size_, other.size_);  
        std::swap(capacity_, other.capacity_);  
    }  
};
```

# Explicitly Deleted Copy & Moves

- We may not want a type to be copyable / moveable.
- Non-copyable types implement “unique ownership” semantics.
  - Only this object owns this resource but can transfer ownership.
- Non-moveable types are rare.
  - Best example is [std::scoped\\_lock](#)
- If we want to delete them, we can declare  
`fn() = delete;`

```
struct T {  
    T(const T&) = delete;  
    T(T&&) = delete;  
  
    T& operator=(const T&) = delete;  
    T& operator=(T&&) = delete;  
};
```

# Implicitly Deleted Copies & Moves

- Under certain conditions, the compiler will not generate the copy and move member functions.
- The implicitly defined copy constructor calls the copy constructor member-wise.
  - If one of its members doesn't have a copy constructor, the compiler can't generate one for you.
  - Same applies for copy assignment, move constructor, and move assignment.
- Under certain conditions, the compiler will not automatically generate copy / move assignment / constructors.
  - If you have manually defined a destructor, the move constructor isn't generated.
  - If you have manually defined a move constructor, the copy constructor isn't generated.
- If you define one of the Five, you should explicitly delete, default, or define all of the Five.
  - If the default behaviour isn't sufficient for one of them, it likely isn't sufficient for others
  - Explicitly doing this tells the reader of your code that you have carefully considered this
  - This also means you don't need to remember all of the rules about "if I write X, then is Y generated?"



# RAII Case Study: Smart Pointers

- Introduced in C++11.
- Ways of wrapping unnamed (i.e., raw pointer) heap objects in named stack objects so that object lifetimes can be managed much easier.
- Supports automatic memory management
  - allocate/deallocate according to RAII.
- Usually two ways of approaching problems:
  - `unique_ptr` + raw pointers
  - `shared_ptr` + `weak_ptr`

Type	Implied Ownership
<code>T*</code>	None
<code>std::unique_ptr&lt;T&gt;</code>	Sole ownership
<code>std::shared_ptr&lt;T&gt;</code>	Shared ownership
<code>std::weak_ptr&lt;T&gt;</code>	None

# std::unique\_ptr

- When the unique pointer is destructed, the underlying object is too.
- Slightly more overhead than raw pointers.
- Can be parameterized with a custom deleter
  - default deleter uses delete.
- Array version exists:
  - Use `std::unique_ptr<T[]>` instead.

```
#include <memory>
#include <iostream>

int main() {
    auto up1 = std::unique_ptr<int>{new int{42}};

    // error: no copy constructor
    auto up2 = up1;

    std::unique_ptr<int> up3;
    // no copy assignment
    up3 = up2;

    up3.reset(up1.release()); // OK
    auto up4 = std::move(up3); // OK

    std::cout << up4.get() << "\n";
    std::cout << *up4 << "\n";
    std::cout << *up1 << "\n";
} // dynamically allocated int freed here.
```

# Raw (Observer) Pointers

- Used to “observe” a `unique_ptr`.
- This is an appropriate use of raw pointers in C++.
- Once the original pointer is destructed, you must ensure you don't access the raw pointers (no checks exist).
- These observers do not have ownership of the pointer.
- Also note the use of `nullptr` in C++ instead of `NULL`.

```
#include <memory>
#include <iostream>

int main() {
    auto up1 = std::unique_ptr<int>{new int{5}};
    auto op1 = up1.get();
    *op1 = 6;

    std::cout << *op1 << "\n"; // prints 6
    std::cout << *up1 << "\n"; // also prints 6

    // free the managed int
    up1.reset(nullptr);

    // undefined behaviour
    std::cout << *op1 << "\n";
}
```

# std::shared\_ptr

- Several objects share ownership of the underlying resource.
- Uses reference counting to avoid premature freeing.
- When a shared pointer is destructed, if it is the only shared pointer left pointing at the object, then the object is destroyed.
- Observers are std::weak\_ptrs instead of raw pointers.
- Like unique\_ptr, an array version is also available.

```
#include <iostream>
#include <memory>

auto main() -> in {
    auto x = std::make_shared<int>(5);

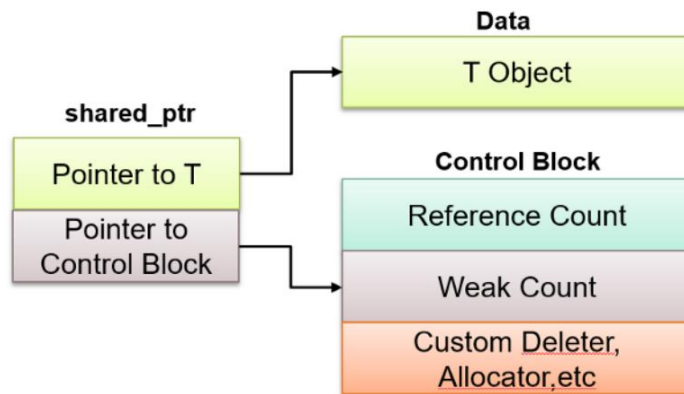
    std::cout << "use count: " << x.use_count() << "\n";
    std::cout << "value: " << *x << "\n";

    x.reset(); // Memory still exists, due to y.
    std::cout << "use count: " << y.use_count() << "\n";
    std::cout << "value: " << *y << "\n";

    // Deletes the memory, since no one else owns the memory
    y.reset();

    std::cout << "use count: " << x.use_count() << "\n";
    std::cout << "value: " << *y << "\n";
}
```

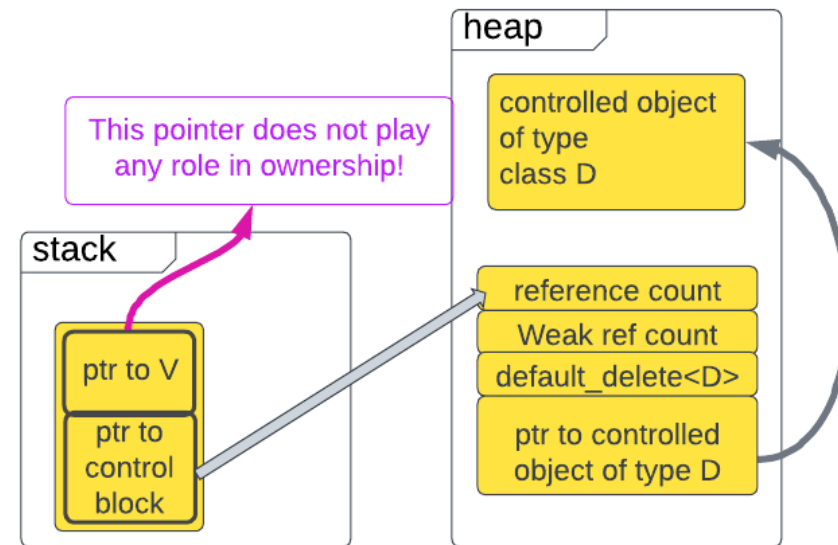
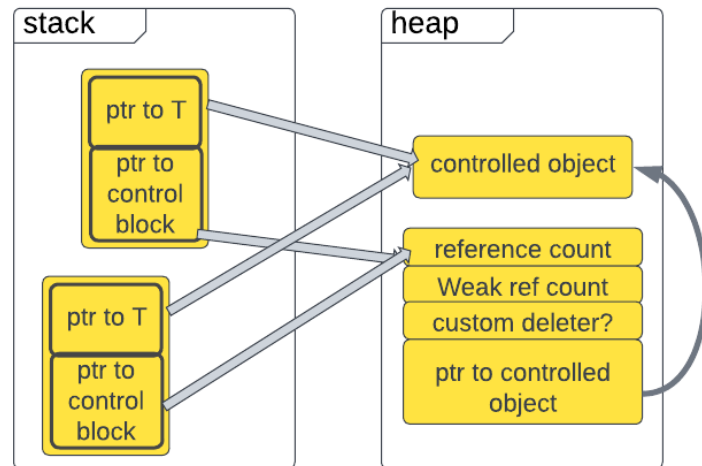
# std::shared\_ptr Control Block



`shared_ptr`, unlike `unique_ptr`, places a layer of indirection between the physical heap-allocated object and the notion of ownership.

`shared_ptr` instances are essentially participating in ref-counted ownership of the control block.

The control block itself is the sole arbiter of what it means to “delete the controlled object.”



# std::weak\_ptr

- Weak pointers are used with share pointers when:
  - You don't want to add to the reference count.
  - You want to be able to check if the underlying data is still valid before using it.
  - Break a circular dependency.
- Must be converted to a shared\_ptr with the lock() method before it can be used.

```
#include <iostream>
#include <memory>

auto main() -> int {
    auto x = std::make_shared<int>(1);

    // x owns the memory
    auto wp = std::weak_ptr<int>(x);

    // now y is a shared_ptr and can be used.
    auto y = wp.lock();

    if (y != nullptr) {
        // x and y own the memory.
        // Do something with y.
        std::cout << "Attempt 1: " << *y << '\n';
    }
}
```

# When to Use Which Pointer

## Unique Pointer vs Shared Pointer

- You almost always want a `unique_ptr` over a `shared_ptr`
- Use a `shared_ptr` if either:

- An object has multiple owners, and you don't know which one will stay around the longest; or
- You need temporary ownership
  - This is very rare.
  - Outside scope of this course.

### “Leak freedom in C++” poster

Strategy	Natural examples	Cost	Rough frequency
1. Prefer <b>scoped lifetime</b> by default (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	O(80%) of objects
2. Else prefer <b>make_unique &amp; unique_ptr</b> or a container, if the object must have its own lifetime (i.e., heap) and ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free <b>Automates</b> simple heap use in a library	
3. Else prefer <b>make_shared &amp; shared_ptr</b> , if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) <b>Automates</b> shared object use in a library	O(20%) of objects

Don't use owning raw `*s` == don't use explicit `delete`

Don't create ownership cycles across modules by owning “upward” (violates layering)  
Use `weak_ptr` to break cycles

Function Signature	Ownership Semantic
<code>func(value)</code>	<ul style="list-style-type: none"><li>▪ Is an independent owner of the resource</li><li>▪ Deletes the resource automatically at the end of <code>func</code></li></ul>
<code>func(pointer*)</code>	<ul style="list-style-type: none"><li>▪ Borrows the resource</li><li>▪ The resource could be empty</li><li>▪ Must not delete the resource</li></ul>
<code>func(reference&amp;)</code>	<ul style="list-style-type: none"><li>▪ Borrows the resource</li><li>▪ The resource could not be empty</li><li>▪ Must not delete the resource</li></ul>
<code>func(std::unique_ptr)</code>	<ul style="list-style-type: none"><li>▪ Is an independent owner of the resource</li><li>▪ Deletes the resource automatically at the end of <code>func</code></li></ul>
<code>func(shared_ptr)</code>	<ul style="list-style-type: none"><li>▪ Is a shared owner of the resource</li><li>▪ May delete the resource at the end of <code>func</code></li></ul>

# Smart Pointer Factory Functions

- `make_shared()` and `make_unique()` wrap `raw new`, just as `~shared_ptr()` and `~unique_ptr()` wrap `raw delete`.
  - Pass the arguments you would have passed to the underlying object's constructor to these functions.
- Never touch raw pointers with hands, and then never need to worry about leaking them.
- `make_unique()` prevents the *unspecified-evaluation-order* leak bug triggered by expressions like:
  - `foo(unique_ptr<T>(new T()), unique_ptr<U>(new U()));`
  - Above, if either the constructor of `T` or `U` throw and the other object has already been allocated, the destructor of `unique_ptr` won't be called.
- `make_shared()` is faster than using `shared_ptr`'s constructors.
  - Able to allocate the managed object *and* control block in a single allocation rather than two.



# Resource Safety

To ensure resource safety in C++, we always attach the lifetime of one object to that of something else.

- **Named objects:**
  - A variable in a function is tied to its scope.
  - A data member is tied to the lifetime of the class instance.
  - An element in a `std::vector` is tied to the lifetime of the vector.
- **Unnamed objects:**
  - A heap object should be tied to the lifetime of whatever object created it.
    - Ideally this should be a smart pointer.
  - Examples of *bad programming practice*:
    - An owning raw pointer is tied to nothing.
    - A heap-allocated C-style array is tied to nothing.

# Feedback (stop recording)

